

```
In [139]: 1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 import plotly.express as px
6 import scipy.stats as stats
7 import statsmodels.api as sm
8 import statsmodels.stats.contingency_tables as ct
9 from statsmodels.graphics.gofplots import qqplot
10 from statsmodels.stats.weightstats import ztest
11 from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error
12 from sklearn.metrics import accuracy_score, classification_report, confusion_matrix, f1_score, recall_score
13 from sklearn.preprocessing import LabelEncoder
14 from sklearn.neural_network import MLPRegressor
15 from sklearn.datasets import make_regression
16 from sklearn.tree import DecisionTreeClassifier
17 from sklearn.model_selection import train_test_split
18 from sklearn.metrics import mean_absolute_error
19 from sklearn.preprocessing import QuantileTransformer
20 from sklearn.ensemble import RandomForestClassifier
21 from sklearn.compose import TransformedTargetRegressor
22 from sklearn.metrics import mean_squared_error
23 from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
24 from sklearn.neighbors import KNeighborsClassifier
25 from sklearn.naive_bayes import GaussianNB
26 from sklearn.tree import DecisionTreeClassifier
27 from sklearn.svm import SVC
28 from sklearn.linear_model import LogisticRegression
29 from sklearn.linear_model import LinearRegression
30 import xgboost as xgb
31 from sklearn.linear_model import Ridge
32 from sklearn.linear_model import Lasso
33 from sklearn.linear_model import ElasticNet
34 from sklearn.preprocessing import PolynomialFeatures
35 from sklearn.model_selection import cross_val_score
36 from sklearn import preprocessing
37 from sklearn.model_selection import KFold
38 from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error
39 from sklearn.model_selection import RandomizedSearchCV
40 from sklearn.metrics import roc_curve, auc
41 from matplotlib.legend_handler import HandlerLine2D
42 from pprint import pprint
43 from utils import *
44 %matplotlib inline
```

```
In [140]: 1 df_train = pd.read_csv('Modeles/trained_Data.csv')
```

```
In [141]: 1 #df_train = df_train.drop(['Unnamed: 0'],axis =1)
2 df_train.loc[df_train['PrixNuitee'] < 100,'classes']=0
3 df_train.loc[(df_train.PrixNuitee >= 100),['classes']]=1
4 #df_train['classes'] = (df['PrixNuitee']>100)
```

```
In [142]: 1 df_train['classes'] = df_train['classes'].astype('int')
```

```
In [143]: 1 Classifiers = {
2     'DecisionTreeClassifier': DecisionTreeClassifier,
3     'RandomForestClassifier': RandomForestClassifier,
4     'KNeighborsClassifier': KNeighborsClassifier,
5     'LogisticRegression': LogisticRegression
6 }
```

```
In [144]: 1 df_train = df_train[['Type_logement','type_propriete','PrixNuitee','NbChambres','Capacite_accueil','Desc_pre','Titre']]
```

```
In [145]: 1 Regressors = {
2     'LinearRegression': LinearRegression,
3     'XGBoostRegression': xgb.XGBRFRegressor,
4     'Ridge': Ridge,
5     'Lasso': Lasso,
6     'Polynomial Regression':PolynomialFeatures
7 }
```

In [146]:

```
1 def classifier(modelclassif,X_class,y_class):
2     #X_class = df_train.drop(['PrixNuıtee','classes'],axis=1)
3     #y_class = df_train['classes']
4     X_train_class, X_test_class, y_train_class, y_test_class = train_test_split(X_class, y_class,random_state=1, tes
5     if modelclassif == 'DecisionTreeClassifier':
6         clf_model = DecisionTreeClassifier(criterion="entropy", random_state=42,
7                                           max_depth=5,
8                                           min_samples_leaf=0.20,
9                                           splitter = 'best'
10                                          )
11     clf_model.fit(X_train_class,y_train_class)
12     y_predict_class = clf_model.predict(X_test_class)
13 elif modelclassif == 'RandomForestClassifier':
14     """
15     {'n_estimators': 400,
16     'min_samples_split': 10,
17     'min_samples_leaf': 4,
18     'max_features': 'auto',
19     'max_depth': 70,
20     'bootstrap': True}
21     """
22     clf_model = RandomForestClassifier(
23         n_estimators = 400,
24         min_samples_split = 10,
25         min_samples_leaf = 4,
26         max_features = 'auto',
27         max_depth = 70,
28         bootstrap = True
29     )
30     clf_model.fit(X_train_class,y_train_class)
31     y_predict_class = clf_model.predict(X_test_class)
32 elif modelclassif == 'KNeighborsClassifier':
33     clf_model = KNeighborsClassifier(n_neighbors = 3)
34     clf_model.fit(X_train_class,y_train_class)
35     y_predict_class = clf_model.predict(X_test_class)
36 elif modelclassif == 'LogisticRegression':
37     clf_model = LogisticRegression()
38     clf_model.fit(X_train_class,y_train_class)
39     y_predict_class = clf_model.predict(X_test_class)
40 acc = accuracy_score(y_test_class,y_predict_class)
41 report = classification_report(y_test_class,y_predict_class)
42 Classifiers[modelclassif] = clf_model
43 return clf_model,acc,report
44
45
46 def regressor(modelReg,X,y):
47     #X_reg = df_train_1.drop(['PrixNuıtee','classes','predictedClass'],axis=1)
48     #y_reg = df_train_1['PrixNuıtee']
49     X_train, X_test, y_train, y_test = train_test_split(X, y,random_state=1, test_size=0.1)
50     if modelReg == 'LinearRegression':
51         reg_model = LinearRegression()
52         reg_model.fit(X_train,y_train)
53         r2 = r2_score(reg_model.predict(X_test),y_test)
54         mse = mean_squared_error(reg_model.predict(X_test),y_test)
55         mae = mean_absolute_error(reg_model.predict(X_test),y_test)
56     elif modelReg == 'XGBoostRegression':
57         reg_model = xgb.XGBRFRegressor()
58         reg_model.fit(X_train,y_train)
59         r2 = r2_score(reg_model.predict(X_test),y_test)
60         mse = mean_squared_error(reg_model.predict(X_test),y_test)
61         mae = mean_absolute_error(reg_model.predict(X_test),y_test)
62     elif modelReg == 'Ridge':
63         reg_model = Ridge()
64         reg_model.fit(X_train,y_train)
65         r2 = r2_score(reg_model.predict(X_test),y_test)
66         mse = mean_squared_error(reg_model.predict(X_test),y_test)
67         mae = mean_absolute_error(reg_model.predict(X_test),y_test)
68     elif modelReg == 'Lasso':
69         reg_model = Lasso()
70         reg_model.fit(X_train,y_train)
71         r2 = r2_score(reg_model.predict(X_test),y_test)
72         mse = mean_squared_error(reg_model.predict(X_test),y_test)
73         mae = mean_absolute_error(reg_model.predict(X_test),y_test)
74     elif modelReg == 'Polynomial Regression':
75         reg_model = LinearRegression()
76         reg_model.fit(X_train,y_train)
77         r2 = r2_score(reg_model.predict(X_test),y_test)
78         mse = mean_squared_error(reg_model.predict(X_test),y_test)
79         mae = mean_absolute_error(reg_model.predict(X_test),y_test)
80     Regressors[modelReg]=reg_model
81     return reg_model, r2,mse,mae
```

DecisionTreeClassifier Parameter tuning :

In [53]:

```
1 criterons = ['gini','entropy']
2 splitters = ['best','random']
3 X_class = df_train.drop(['PrixNuitee','classes'],axis=1)
4 y_class = df_train['classes']
5 X_train_class, X_test_class, y_train_class, y_test_class = train_test_split(X_class, y_class,random_state=1, test_size=0.2)
6
7 x_train = X_train_class
8 y_train = y_train_class
9 x_test = X_test_class
10 y_test = y_test_class
```

criterion

In [52]:

```
1 def criterionP(c):
2     clf_model = DecisionTreeClassifier(criterion=c, random_state=42)
3     clf_model.fit(X_train_class,y_train_class)
4     return accuracy_score(clf_model.predict(X_test_class),y_test_class)
5
6 #criterionP('gini')
7
8 for c in criterons:
9     print(c, ' : ',criterionP(c))
```

```
gini : 0.7938931297709924
entropy : 0.7900763358778626
```

splitter

In [54]:


```
1 def splitterP(s):
2     clf_model = DecisionTreeClassifier(splitter=s, random_state=42)
3     clf_model.fit(X_train_class,y_train_class)
4     return accuracy_score(clf_model.predict(X_test_class),y_test_class)
5
6 #criterionP('gini')
7
8 for s in splitters:
9     print(s, ' : ',splitterP(s))
```

```
best : 0.7938931297709924
random : 0.7919847328244275
```

max_depth

In [59]:


```
1 max_depths = np.linspace(1, 32, 32, endpoint=True)
2 train_results = []
3 test_results = []
4 for max_depth in max_depths:
5     dt = DecisionTreeClassifier(max_depth=max_depth)
6     dt.fit(X_train_class, y_train_class)
7     train_pred = dt.predict(X_train_class)
8     false_positive_rate, true_positive_rate, thresholds = roc_curve(y_train_class, train_pred)
9     roc_auc = auc(false_positive_rate, true_positive_rate)
10    # Add auc score to previous train results
11    train_results.append(roc_auc)
12    y_pred = dt.predict(X_test_class)
13    false_positive_rate, true_positive_rate, thresholds = roc_curve(y_test_class, y_pred)
14    roc_auc = auc(false_positive_rate, true_positive_rate)
15    # Add auc score to previous test results
16    test_results.append(roc_auc)
17
18 line1, = plt.plot(max_depths, train_results, 'b', label="Train AUC")
19 line2, = plt.plot(max_depths, test_results, 'r', label="Test AUC")
20 plt.legend(handler_map={line1: HandlerLine2D(numpoints=2)})
21 plt.ylabel('AUC score')
22 plt.xlabel('Tree depth')
23 plt.show()
```



min_samples_split:

In [61]:


```
1 min_samples_splits = np.linspace(0.1, 1.0, 10, endpoint=True)
2 train_results = []
3 test_results = []
4
5 for min_samples_split in min_samples_splits:
6     dt = DecisionTreeClassifier(min_samples_split=min_samples_split)
7     dt.fit(x_train, y_train)
8     train_pred = dt.predict(x_train)
9     false_positive_rate, true_positive_rate, thresholds = roc_curve(y_train, train_pred)
10    roc_auc = auc(false_positive_rate, true_positive_rate)
11    train_results.append(roc_auc)
12    y_pred = dt.predict(x_test)
13    false_positive_rate, true_positive_rate, thresholds = roc_curve(y_test, y_pred)
14    roc_auc = auc(false_positive_rate, true_positive_rate)
15    test_results.append(roc_auc)
16
17
18 line1, = plt.plot(min_samples_splits, train_results, 'b', label="Train AUC")
19 line2, = plt.plot(min_samples_splits, test_results, 'r', label="Test AUC")
20 plt.legend(handler_map={line1: HandlerLine2D(numpoints=2)})
21 plt.ylabel('AUC score')
22 plt.xlabel('min samples split')
23 plt.show()
24
25 # Min samples -> 0.7
```



min samples leafs

In [62]:

```
1 min_samples_leafs = np.linspace(0.1, 0.5, 5, endpoint=True)
2 train_results = []
3 test_results = []
4 for min_samples_leaf in min_samples_leafs:
5     dt = DecisionTreeClassifier(min_samples_leaf=min_samples_leaf)
6     dt.fit(x_train, y_train)
7     train_pred = dt.predict(x_train)
8     false_positive_rate, true_positive_rate, thresholds = roc_curve(y_train, train_pred)
9     roc_auc = auc(false_positive_rate, true_positive_rate)
10    train_results.append(roc_auc)
11    y_pred = dt.predict(x_test)
12    false_positive_rate, true_positive_rate, thresholds = roc_curve(y_test, y_pred)
13    roc_auc = auc(false_positive_rate, true_positive_rate)
14    test_results.append(roc_auc)
15
16 line1, = plt.plot(min_samples_leafs, train_results, 'b', label="Train AUC")
17 line2, = plt.plot(min_samples_leafs, test_results, 'r', label="Test AUC")
18 plt.legend(handler_map={line1: HandlerLine2D(numpoints=2)})
19 plt.ylabel('AUC score')
20 plt.xlabel('min samples leafs')
21 plt.show()
22
23 # Min samples leafs --> 0.20
```



RandomForestClassifier Parameter tuning :

```
In [66]: 1 # Number of trees in random forest
2 n_estimators = [int(x) for x in np.linspace(start = 200, stop = 2000, num = 10)]
3 # Number of features to consider at every split
4 max_features = ['auto', 'sqrt']
5 # Maximum number of levels in tree
6 max_depth = [int(x) for x in np.linspace(10, 110, num = 11)]
7 max_depth.append(None)
8 # Minimum number of samples required to split a node
9 min_samples_split = [2, 5, 10]
10 # Minimum number of samples required at each leaf node
11 min_samples_leaf = [1, 2, 4]
12 # Method of selecting samples for training each tree
13 bootstrap = [True, False]
14
15 # Create the random grid
16 random_grid = {'n_estimators': n_estimators,
17                 'max_features': max_features,
18                 'max_depth': max_depth,
19                 'min_samples_split': min_samples_split,
20                 'min_samples_leaf': min_samples_leaf,
21                 'bootstrap': bootstrap}
22
23 pprint(random_grid)
```

```
{'bootstrap': [True, False],
 'max_depth': [10, 20, 30, 40, 50, 60, 70, 80, 90, 100, 110, None],
 'max_features': ['auto', 'sqrt'],
 'min_samples_leaf': [1, 2, 4],
 'min_samples_split': [2, 5, 10],
 'n_estimators': [200, 400, 600, 800, 1000, 1200, 1400, 1600, 1800, 2000]}
```

```
In [71]: 1 # Use the random grid to search for best hyperparameters
2 # First create the base model to tune
3 rf = RandomForestClassifier(random_state = 42)
4
5
6 rf_random = RandomizedSearchCV(estimator=rf, param_distributions=random_grid,
7                                n_iter = 100, scoring='accuracy', #'f1'
8                                cv = 3, verbose=2, random_state=42, n_jobs=-1,
9                                return_train_score=True)
10
11 # Fit the random search model
12 rf_random.fit(X_train_class, y_train_class);
```

Fitting 3 folds for each of 100 candidates, totalling 300 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 4 concurrent workers.
[Parallel(n_jobs=-1)]: Done 33 tasks      | elapsed: 1.6min
[Parallel(n_jobs=-1)]: Done 154 tasks    | elapsed: 6.8min
[Parallel(n_jobs=-1)]: Done 300 out of 300 | elapsed: 13.4min finished
```

```
In [138]: 1 X_class
```

...


```
In [72]: 1 rf_random.best_params_
```

```
Out[72]: {'n_estimators': 400,
 'min_samples_split': 10,
 'min_samples_leaf': 4,
 'max_features': 'auto',
 'max_depth': 70,
 'bootstrap': True}
```

KNeighborsClassifier Parameter tuning :

n_neighbours

```
In [80]: 1 neighbors = list(range(1,30))
2 train_results = []
3 test_results = []
4 for n in neighbors:
5     model = KNeighborsClassifier(n_neighbors=n)
6     model.fit(x_train, y_train)
7     train_pred = model.predict(x_train)
8     false_positive_rate, true_positive_rate, thresholds = roc_curve(y_train, train_pred)
9     roc_auc = auc(false_positive_rate, true_positive_rate)
10    train_results.append(roc_auc)
11    y_pred = model.predict(x_test)
12    false_positive_rate, true_positive_rate, thresholds = roc_curve(y_test, y_pred)
13    roc_auc = auc(false_positive_rate, true_positive_rate)
14    test_results.append(roc_auc)
15
16
17 line1, = plt.plot(neighbors, train_results, 'b', label="Train AUC")
18 line2, = plt.plot(neighbors, test_results, 'r', label="Test AUC")
19 plt.legend(handler_map={line1: HandlerLine2D(numpoints=2)})
20 plt.ylabel('AUC score')
21 plt.xlabel('n_neighbors')
22 plt.show()
23
24 # n_neighbours ----> 3
```



Regressors :

```
In [84]: 1 X_class = df_train.drop(['PrixNuitee', 'classes'],axis=1)
2 y_class = df_train['PrixNuitee']
3 X_train_class, X_test_class, y_train_class, y_test_class = train_test_split(X_class, y_class,random_state=1, test_s:
4
5 x_train = X_train_class
6 y_train = y_train_class
7 x_test = X_test_class
8 y_test = y_test_class
```

XGBoostRegression Tuning

```
In [175]: 1 params = { 'max_depth': [1,4,5,10],
2             'learning_rate': [0.01, 0.05, 0.1],
3             'n_estimators': [100, 500, 1000],
4             'colsample_bytree': [0.3, 0.7,0.8]}
5
6 xgbr = xgb.XGBRegressor(seed = 10)
7 pprint(params)
```

```
{'colsample_bytree': [0.3, 0.7, 0.8],
'learning_rate': [0.01, 0.05, 0.1],
'max_depth': [1, 4, 5, 10],
'n_estimators': [100, 500, 1000]}
```

```
In [176]: 1 rf_random = RandomizedSearchCV(estimator=xgbr, param_distributions=params,
2                                 n_iter = 100, scoring='r2', #'f1'
3                                 cv = 3, verbose=2, random_state=42, n_jobs=-1,
4                                 return_train_score=True)
5
6 # Fit the random search model
7 rf_random.fit(X_train_class, y_train_class);
```

...

```
In [ ]: 1 rf_random
```

Model Training

Classifiers :


```
In [147]: 1 X_class = df_train.drop(['PrixNuitee', 'classes'],axis=1)
          2 y_class = df_train['classes']
```

```
In [178]: 1 X_class
```

Out[178]:

	Type_logement	type_propriete	NbChambres	Capacite_accueil	Desc_pre	Titre_pre	Reg_pre
0	1	9	1	2	547.415747	1.000000	70.242437
1	2	0	1	2	1656.273528	2.645751	54.027771
2	1	9	1	2	1204.822394	3.162278	52.009614
3	2	0	2	4	3745.453377	0.000000	136.319478
4	2	0	1	4	2751.334404	4.123106	70.121323
...
5229	2	0	1	4	4637.633233	36.013886	146.003425
5230	1	0	1	2	6147.921275	128.432862	94.090382
5231	1	0	1	3	17006.853912	104.033648	759.368158
5232	2	9	1	4	5577.381016	105.792249	761.626549
5233	2	7	2	5	5746.660334	98.848369	60.049979

5234 rows × 7 columns

```
In [148]: 1 for model in Classifiers.keys():
          2     _,acc,report = classifier(model,X_class,y_class)
          3     print(model,' accuracy => ',acc)
          4     print(report)
```

DecisionTreeClassifier accuracy => 0.833969465648855

	precision	recall	f1-score	support
0	0.92	0.87	0.89	426
1	0.54	0.68	0.61	98
accuracy			0.83	524
macro avg	0.73	0.78	0.75	524
weighted avg	0.85	0.83	0.84	524

RandomForestClassifier accuracy => 0.8645038167938931

	precision	recall	f1-score	support
0	0.89	0.95	0.92	426
1	0.70	0.48	0.57	98
accuracy			0.86	524
macro avg	0.79	0.72	0.74	524
weighted avg	0.85	0.86	0.85	524

KNeighborsClassifier accuracy => 0.7748091603053435

	precision	recall	f1-score	support
0	0.83	0.91	0.87	426
1	0.31	0.17	0.22	98
accuracy			0.77	524
macro avg	0.57	0.54	0.55	524
weighted avg	0.73	0.77	0.75	524

LogisticRegression accuracy => 0.8492366412213741

	precision	recall	f1-score	support
0	0.85	0.99	0.91	426
1	0.83	0.24	0.38	98
accuracy			0.85	524
macro avg	0.84	0.62	0.65	524
weighted avg	0.85	0.85	0.81	524

/home/mouad/.virtualenvs/ml/lib/python3.8/site-packages/sklearn/linear_model/_logistic.py:762: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)
n_iter_i = _check_optimize_result(

Best Classifier :

```
In [149]: 1 ## Python program to understand the usage of tabulate function for printing tables in a tabular format
2 from tabulate import tabulate
3 data = [['Accuracy']+accuracy_score(Classifiers[key].predict(X_class),y_class) for key in Classifiers.keys()],
4         ['F1']+f1_score(Classifiers[key].predict(X_class),y_class) for key in Classifiers.keys()],
5         ['Recall']+recall_score(Classifiers[key].predict(X_class),y_class) for key in Classifiers.keys()]
6         ]
7 print(tabulate(data, headers=['Measure']+list(Classifiers.keys())))
```

Measure	DecisionTreeClassifier	RandomForestClassifier	KNeighborsClassifier	LogisticRegression
Accuracy	0.823271	0.900459	0.837982	0.831295
F1	0.624442	0.712314	0.485437	0.366116
Recall	0.548111	0.858855	0.680272	0.765766

```
In [180]: 1 import joblib
2
3 # save the model to disk
4 filename = 'Result/class_model.sav'
5 joblib.dump(Classifiers['RandomForestClassifier'], filename)
6
7 # some time later...
8
9 # load the model from disk
10 loaded_model = joblib.load(filename)
```

Regressors :

RegressorOne :

```
In [150]: 1 df_train_1 = df_train.loc[df_train['classes']== 0]
```

```
In [151]: 1 X_reg_1 = df_train_1.drop(['PrixNuitee','classes'],axis=1)
2 y_reg_1 = df_train_1['PrixNuitee']
```

```
In [152]: 1 for model in Regressors.keys():
2     _,r2,mse,mae = regressor(model,X_reg_1,y_reg_1)
3     print(model, ' -----\n',
4           'r2 => ',r2,
5           '\nMean Squared Error => ',mse,
6           '\nMean Absolute Error => ',mae
7           )
8
9 regressorOneRegister = Regressors
```

LinearRegression -----
r2 => -1.0896314951299413
Mean Squared Error => 286.21439911886216
Mean Absolute Error => 13.308734555898738
XGBoostRegression -----
r2 => -0.7841473199379532
Mean Squared Error => 276.13484920558153
Mean Absolute Error => 12.93759408978184
Ridge -----
r2 => -1.091913298594402
Mean Squared Error => 286.2023085209517
Mean Absolute Error => 13.309329820568477
Lasso -----
r2 => -2.0277585861780976
Mean Squared Error => 290.8839475060331
Mean Absolute Error => 13.70147321778131
Polynomial Regression -----
r2 => -1.0896314951299413
Mean Squared Error => 286.21439911886216
Mean Absolute Error => 13.308734555898738

Best Regressor I

```
In [153]: 1 data = [['R2']+r2_score(regressorOneRegister[key].predict(X_reg_1),y_reg_1) for key in regressorOneRegister.keys()],
2         ['MSE']+mean_squared_error(regressorOneRegister[key].predict(X_reg_1),y_reg_1) for key in regressorOneRegister.keys()],
3         ['MAE']+mean_absolute_error(regressorOneRegister[key].predict(X_reg_1),y_reg_1) for key in regressorOneRegister.keys()]
4         ]
5 print(tabulate(data, headers=['Measure']+list(regressorOneRegister.keys())))
```

Measure	LinearRegression	XGBoostRegression	Ridge	Lasso	Polynomial Regression
R2	-0.664963	-0.349655	-0.66694	-1.49731	-0.664963
MSE	245.415	213.913	245.414	253.849	245.415
MAE	12.4531	11.6332	12.4536	12.8354	12.4531


```
In [181]: 1 import joblib
2
3 # save the model to disk
4 filename = 'Result/regressOne_model.sav'
5 joblib.dump(regressorOneRegister['XGBoostRegression'], filename)
6
7 # some time later...
8
9 # load the model from disk
10 loaded_model = joblib.load(filename)
```

RegressorTwo :

```
In [154]: 1 df_train_2 = df_train.loc[df_train['classes']== 1]
2 X_reg_2 = df_train_2.drop(['PrixNuitee', 'classes'],axis=1)
3 y_reg_2 = df_train_2['PrixNuitee']
```

```
In [155]: 1 for model in Regressors.keys():
2     _,r2,mse,mae = regressor(model,X_reg_2,y_reg_2)
3     print(model,' -----\n',
4           'r2 => ',r2,
5           '\nMean Squared Error => ',mse,
6           '\nMean Absolute Error => ',mae
7         )
8
9 regressorTwoRegister = Regressors
```

LinearRegression -----
r2 => -2.077629148949636
Mean Squared Error => 8427.567690417067
Mean Absolute Error => 52.31301014547509
XGBoostRegression -----
r2 => -0.19234252105246696
Mean Squared Error => 7065.416932234301
Mean Absolute Error => 52.60723747397369
Ridge -----
r2 => -2.09602616573797
Mean Squared Error => 8438.8793151152
Mean Absolute Error => 52.307518872797175
Lasso -----
r2 => -2.561410922667501
Mean Squared Error => 8750.610776040523
Mean Absolute Error => 52.52713463270967
Polynomial Regression -----
r2 => -2.077629148949636
Mean Squared Error => 8427.567690417067
Mean Absolute Error => 52.31301014547509

Best Regressor II

```
In [156]: 1 data = [['R2']+r2_score(regressorTwoRegister[key].predict(X_reg_1),y_reg_1) for key in regressorTwoRegister.keys()]
2         ['MSE']+mean_squared_error(regressorTwoRegister[key].predict(X_reg_1),y_reg_1) for key in regressorTwoRegis
3         ['MAE']+mean_absolute_error(regressorTwoRegister[key].predict(X_reg_1),y_reg_1) for key in regressorTwoRegi
4         ]
5 print(tabulate(data, headers=['Measure']+list(regressorTwoRegister.keys())))
```

Measure	LinearRegression	XGBoostRegression	Ridge	Lasso	Polynomial Regression
R2	-6.82641	-20.9524	-6.96151	-9.92741	-6.82641
MSE	6557.12	8132.25	6503.97	5689.37	6557.12
MAE	73.0202	85.4656	72.8571	70.3319	73.0202

Model Predicting:

In [170]:

```
1 def predict(X):
2     #----- Classifier
3     clf_model = Classifiers['RandomForestClassifier']#BEST CLASSIFIER
4     classes = clf_model.predict(X.drop(['PrixNuitee', 'classes'],axis=1))
5     X['classes']=classes
6     #----- Regressor I
7     X_1 = X.loc[X['classes']== 0]
8     y_1 = X_1['PrixNuitee'].values
9     reg_modelI = regressorOneRegister['XGBoostRegression']#BEST REGRESOR|ONE
10    prixNuitee_1 = reg_modelI.predict(X_1.drop(['PrixNuitee', 'classes'],axis=1))
11    #----- Regressor II
12    X_2 = X.loc[X['classes']== 1]
13    reg_modelII = regressorTwoRegister['XGBoostRegression']#BEST REGRESSOR|TWO
14    y_2 = X_2['PrixNuitee'].values
15    prixNuitee_2 = reg_modelII.predict(X_2.drop(['PrixNuitee', 'classes'],axis=1))
16    return prixNuitee_1,y_1,prixNuitee_2,y_2
17
18 def predict_Adv(X):
19     #[['Longitude','Latitude','type_propriete','PrixNuitee',
20     #'NbChambres','Capacite_accueil','Desc_pre','Titre_pre','Reg_pre']]
21     #----- Classifier
22     clf_model = Classifiers['RandomForestClassifier']#BEST CLASSIFIER
23     classes = clf_model.predict(X.drop(['latitude', 'longitude']))
24     X['classes']=classes
25     #----- Regressor I
26     X_1 = X.loc[X['classes']== 0]
27     reg_modelI = regressorOneRegister['XGBoostRegression']#BEST REGRESOR|ONE
28     prixNuitee_1 = reg_modelI.predict(X_1.drop(['latitude', 'longitude'],axis=1))
29     result_1 = pd.DataFrame(data = {
30         'latitude': X_1['latitude'].values,
31         'longitude':X_1['longitude'].values,
32         'PrixNuitee':prixNuitee_1
33     })
34     )
35     #----- Regressor II
36     X_2 = X.loc[X['classes']== 1]
37     reg_modelII = regressorTwoRegister['XGBoostRegression']#BEST REGRESSOR|TWO
38     prixNuitee_2 = reg_modelII.predict(X_2.drop(['latitude', 'longitude'],axis=1))
39     result_2 = pd.DataFrame( data = {
40         'latitude':X_2['latitude'].values,
41         'longitude':X_2['longitude'].values,
42         'PrixNuitee':prixNuitee_2}
43     )
44     #----- File Json
45     result = pd.concat([result_1, result_2], ignore_index=True)
46     result.to_json('predict.json',orient='records')
47     return result #prixNuitee_1,prixNuitee_2
48
```

In [182]:

```
1 import joblib
2
3 # save the model to disk
4 filename = 'Result/regressTwo_model.sav'
5 joblib.dump(regressorTwoRegister['XGBoostRegression'], filename)
6
7 # some time later...
8
9 # load the model from disk
10 loaded_model = joblib.load(filename)
```

In []:

```
1
```

In [171]:

```
1 prixNuitee_1,y_1,prixNuitee_2,y_2 = predict(df_train)
```

In [183]:

```
1 df_train
```

Out[183]:

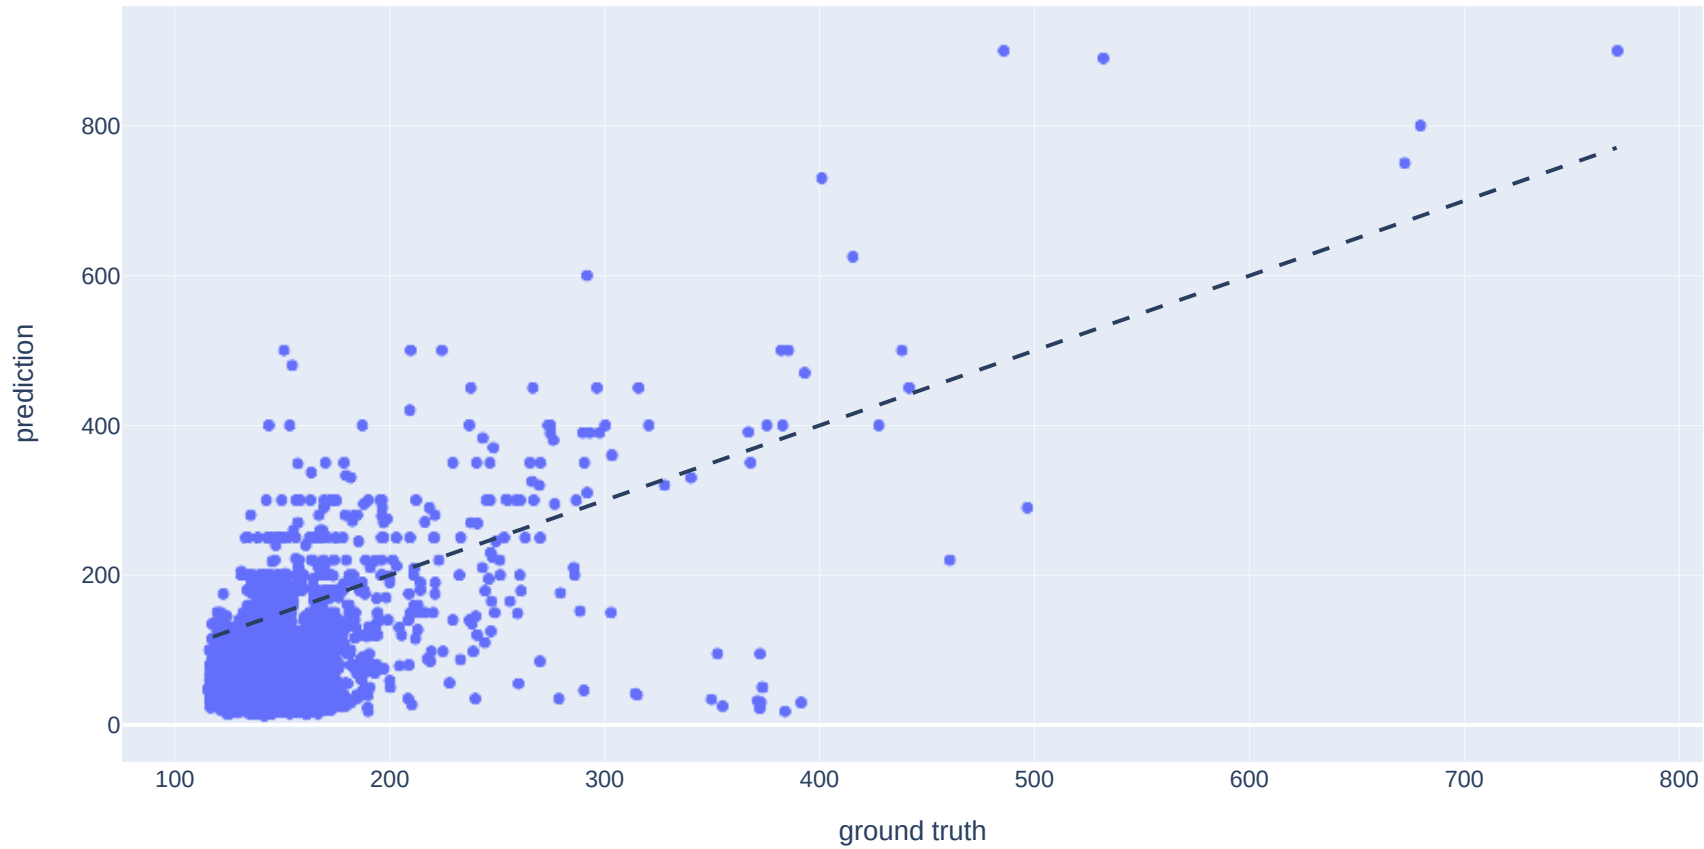
	Type_logement	type_propriete	PrixNuitee	NbChambres	Capacite_accueil	Desc_pre	Titre_pre	Reg_pre	classes
0	1	9	25	1	2	547.415747	1.000000	70.242437	0
1	2	0	71	1	2	1656.273528	2.645751	54.027771	0
2	1	9	75	1	2	1204.822394	3.162278	52.009614	0
3	2	0	155	2	4	3745.453377	0.000000	136.319478	0
4	2	0	80	1	4	2751.334404	4.123106	70.121323	0
...
5229	2	0	60	1	4	4637.633233	36.013886	146.003425	0
5230	1	0	50	1	2	6147.921275	128.432862	94.090382	0
5231	1	0	20	1	3	17006.853912	104.033648	759.368158	0
5232	2	9	70	1	4	5577.381016	105.792249	761.626549	0
5233	2	7	104	2	5	5746.660334	98.848369	60.049979	1

5234 rows × 9 columns

In [172]:

```
1 prixNuitee = list(prixNuitee_1)+list(prixNuitee_2)
2 y_ =list(y_1)+list(y_2)
```

```
In [174]: 1 fig = px.scatter(x=prixNuitee, y=y_, labels={'x': 'ground truth', 'y': 'prediction'})
2 fig.add_shape(
3     type="line", line=dict(dash='dash'),
4     x0=prixNuitee_2.min(), y0=prixNuitee_2.min(),
5     x1=prixNuitee_2.max(), y1=prixNuitee_2.max()
6 )
7 fig.show()
```



```
In [ ]: 1 predict_Adv()
```

```
In [184]: 1 df_train
```

Out[184]:

	Type_logement	type_propriete	PrixNuitee	NbChambres	Capacite_accueil	Desc_pre	Titre_pre	Reg_pre	classes
0	1	9	25	1	2	547.415747	1.000000	70.242437	0
1	2	0	71	1	2	1656.273528	2.645751	54.027771	0
2	1	9	75	1	2	1204.822394	3.162278	52.009614	0
3	2	0	155	2	4	3745.453377	0.000000	136.319478	0
4	2	0	80	1	4	2751.334404	4.123106	70.121323	0
...
5229	2	0	60	1	4	4637.633233	36.013886	146.003425	0
5230	1	0	50	1	2	6147.921275	128.432862	94.090382	0
5231	1	0	20	1	3	17006.853912	104.033648	759.368158	0
5232	2	9	70	1	4	5577.381016	105.792249	761.626549	0
5233	2	7	104	2	5	5746.660334	98.848369	60.049979	1

5234 rows × 9 columns

```
In [ ]: 1
```